Exploring Workplace Mental Health through Data Science: Insights for Organizational Well-Being and Policy Design

Arkadyuti Bhattacharjee (2025MS006)
Ushosree Raha (2025MS027)
Yazhini K (2025MS032)

Under the guidance of Dr. Yashwant Singh Patel

Atal Bihari Vajpayee - Indian Institute of Information Technology and Management (ABV-IIITM), Gwalior

10 November 2025

Contents

List of Figures

Abstract

This report presents a comprehensive, practical and research-informed analysis of mental health in the workplace using the Kaggle "Mental Health in Technology Survey" dataset. The work is structured around three pillars: (1) careful data preparation and responsible analytics, (2) predictive modeling using robust ensemble techniques culminating in a stacking classifier, and (3) visualization and policy simulation supporting evidence-based decisions. The heart of the deliverable is an Early Warning Index (EWI), a straightforward numeric score derived from the models that helps organizations prioritize interventions. The EWI is visualized through distribution plots, demographic comparisons, global mapping, and KPI indicators so that non-technical stakeholders can interpret trends quickly and clearly.

The machine learning work includes Random Forest and XGBoost baseline models and a Stacking Ensemble (XGBoost + Random Forest with Logistic Regression as the meta-learner). After thorough cross-validation and hyperparameter tuning, the stacking model achieved the best test accuracy (approximately 0.6825). However, the report emphasizes interpretability and responsible use: we present feature importance, confusion matrices, and classification reports so that stakeholders can see both overall accuracy and the class-level trade-offs such as false negatives and false positives. This helps organizations decide how to tune the sensitivity for early detection vs. specificity to reduce false alarms.

Beyond technical results, the report simulates a policy intervention — a 20% reduction in the "work interference"— and demonstrates its impact on average EWI in the population. The simulation shows measurable reductions in risk, indicating that modest workplace reforms (e.g., flexible hours, manager training, workload rebalancing) can produce meaningful improvements in employee wellbeing. To complement the quantitative analysis, we include practical policy recommendations oriented to HR teams and government agencies, plus a roadmap for the deployment of a privacy-first EWI dashboard.

1. Introduction

1.1 Background

Mental health is increasingly recognized as a crucial component of employee well-being and organizational performance. Over the last decade, both the academic and policy communities have made sustained efforts to quantify the economic and social burden of poor mental health. For example, global estimates indicate that depression and anxiety disorders cause hundreds of billions of dollars in lost productivity each year. Organizations now face the dual challenge of supporting employees ethically while also maintaining operational performance. Given the sensitivity of mental health data and the often-noted stigma in many workplaces, any analytical approach must balance predictive power with privacy, fairness, and interpretability.

Organizations in high-pressure industries — including technology — encounter unique stressors such as long working hours, performance-based assessments, rapid delivery cycles, and high cognitive load. These environmental factors interact with personal characteristics such as family history and prior treatment, rendering mental health a multifactorial problem. Consequently, static measures—such as an annual survey—offer a limited snapshot. Continuous, responsibly managed analytics can provide rolling insights into emerging trends, enabling earlier support and a more targeted allocation of resources. The Early Warning Index (EWI) introduced in this project aims to be such a tool: simple to interpret, grounded in model outputs, and designed for aggregated, anonymized monitoring rather than intrusive individual surveillance.

Despite the promise of data-driven detection, practical challenges remain. Data quality varies across organizations and surveys, missingness can be non-random, and social desirability bias may influence responses. There is also the legitimate concern of misuse: an employer might react punitively if individual-level risk is disclosed. Therefore, the project's guiding principle is "aggregate and act" — models to detect population or group-level risk trends and design policies that improve workplace conditions, rather than provide disciplinary evidence. This framing ensures compliance with ethical norms, GDPR-like data protection expectations, and long-term trust-building with employees.

Finally, the convergence of academic research, corporate responsibility, and public policy creates an opportunity to leverage machine learning for the public good. Government agencies can use aggregated EWI metrics to inform labor regulations and public health outreach, while organizations can use the same indicators to prioritize wellness investments. The rest of this chapter outlines the problem statement and objectives that motivated this project and explains why a combined research-and-policy approach offers the most realistic path to impact.

1.2 Problem Statement

Workplace mental health is a complex, multi-layered problem shaped by personal, interpersonal, and institutional factors. Employees may experience stressors related to workload, unclear role expectations, insufficient benefits, managerial styles, or personal life pressures. Many of these drivers do not exist in isolation but interact, producing non-linear risk patterns. Traditional HR interventions often rely on point-in-time surveys or reactive policies, which limit the ability to anticipate crises. This project asks whether routinely collected survey data can be harnessed to create an early-warning capability that detects rising group-level risk before it culminates in performance drops or severe individual crises.

Another aspect of the problem is measurement and trust. Employees may fear breaching privacy or exposing vulnerabilities; as a result, direct self-report measures underestimate the actual prevalence. Analytical frameworks must therefore be transparent, explainable, and focused on systemic improvements. The challenge is to construct models that are accurate enough to be useful, yet interpretable enough for HR professionals and policymakers to trust and act on. In addition, any operational deployment must protect employee anonymity, use aggregated insights for decision-making, and incorporate manual oversight to prevent automated or discriminatory actions.

A further dimension is policy translation: even when indicators exist, there is often a gap between evidence and policy implementation. Employers may recognize risk but lack the resources, incentives, or technical capability to act. Thus, the project includes policy simulations — using model outputs to estimate how specific interventions could plausibly change outcomes at scale. Estimating potential impact is essential for producing policy-ready recommendations that can be prioritized based on expected benefit and cost.

Finally, this project seeks to demonstrate a replicable pipeline that organizations can adopt. Instead of a one-off analysis, the goal is to provide a workflow—from data cleaning to modeling to dashboarding—that can be operationalized, audited, and improved over time. The emphasis on reproducibility and a clear governance model makes the outcomes more likely to be sustainable and ethically implemented.

1.3 Objectives

The objectives of this project are technical, practical, and policy-focused. Technically, we aim to build robust predictive models that can distinguish higher-risk respondents from lower-risk ones using relevant survey features. Practically, we aim to consolidate the model outputs into a single, interpretable Early Warning Index (EWI) and a suite of visualizations suited for HR and executive audiences. Policy-wise, the project will quantify hypothetical interventions — for example, reducing work interference — to demonstrate plausible gains from targeted policy measures.

Specifically, the project will:

1. Clean and preprocess the Kaggle survey dataset and create a production-friendly data pipeline.
2. Train Random Forest and XGBoost models and combine them in a Stacking Ensemble to improve predictive performance.
3. Evaluate models with stratified train-test splits and robust cross-validation, reporting accuracy and class-level metrics.
4. Create an Early Warning Index and visual dashboards (EWI distribution, EWI by demographic, global EWI map, KPI indicators).
5. Simulate policy scenarios and translate results into prioritized recommendations for employers and policymakers.

Each of these objectives balances methodological rigor with practical constraints, such as data limitations and privacy considerations. The rest of the document describes how these objectives were met and provides an interpretation of results for both technical and non-technical audiences.

2. Literature Review

The literature on workplace mental health emphasizes both human and economic costs. Several global reports, including those from the World Health Organization and the ILO, quantify the losses from untreated mental health conditions in terms of reduced productivity, increased absenteeism, and higher turnover. Studies show that untreated mental illness is associated with lower job retention and greater healthcare utilization, which together impose tangible costs on both firms and public health systems. There is growing consensus that preventive workplace interventions—ranging from flexible schedules to manager training programs—are cost-effective in many corporate contexts. Empirical studies, often industry-specific, suggest that early interventions yield positive returns through improved morale and reduced sick leave, although the size of the effect can vary by country, sector, and labor market conditions.

Methodological research in occupational health emphasizes the importance of combining individual and contextual factors to explain mental health outcomes. For example, ecological models that incorporate personal history, workplace culture, and macroeconomic stressors provide richer explanations than individual-level models alone. Survey studies often capture these layers through items measuring managerial support, remote-work arrangements, and perceived stigma. Complementary research leverages administrative records and digital traces (e.g., email metadata, calendar density) to detect stress patterns, though these approaches raise more severe privacy questions. The convergence of survey-based and behavioral data suggests a hybrid approach where sensitive models prefer aggregate outputs, aligning with the ethical constraints discussed in privacy literature.

Machine learning has increasingly been applied in mental health research as a tool for pattern detection rather than purely for black-box prediction. Random Forests and gradient-boosting machines (like XGBoost) are popular because they handle heterogeneous and sparse data well and provide useful measures such as feature importances. Stacking ensembles further improve generalization by combining complementary learners. Yet literature cautions that predictive performance must be balanced with interpretability — particularly in domains affecting human welfare. Methods such as SHAP values, partial dependence plots, and simple one-page model summaries have been recommended in numerous applied papers to ensure that stakeholders can make sense of model outputs and trust their implications.

Policy research draws upon these technical findings to recommend governance frameworks for workplace mental health that connect measurement to action. Successful programs combine monitoring (surveys, KPIs), interventions (manager training, access to care), and evaluation (pre-post measures). There is mounting evidence that public-private partnerships and financial incentives (tax credits, subsidies) can accelerate adoption, especially among small and medium enterprises that lack in-house resources. Cross-national studies highlight that policy effectiveness depends on social safety nets and healthcare infrastructure, which influence how easily employees can access mental health services once identified.

Ethics and privacy form a substantial body of work relevant to any applied analytics project in mental health. Guidelines from professional associations and data-protection regulators emphasize consent, minimization, and transparency. The literature recommends storing and reporting results at aggregated levels, using anonymization and differential privacy techniques where needed, and implementing robust governance to control who can access sensitive outputs. This project follows those recommendations by designing an EWI intended for group-level monitoring, by proposing clear operational rules, and by suggesting privacy-preserving options for future production deployments.

3. Resources and Data

3.1 Dataset Description

The analysis uses the Kaggle dataset named "Mental Health in Tech Survey," which collects self-reported information from technology-sector professionals. The dataset includes demographic fields (age, gender, country), employment context (company size, remote work status, tech company indicator), mental-health-related items (history of treatment, access to benefits, work interference), and free-text comments. The dataset is not a representative population sample but reflects the experiences and reporting behavior of tech-sector respondents who volunteered for the survey. Despite potential sampling limitations, the dataset is rich in variables relevant to workplace mental health and therefore well suited for exploratory analysis and methodology demonstration.

Data cleaning is an essential first step: standardizing column names, handling missing values, and converting categorical variables into suitable formats (one-hot encoding or ordinal mapping as appropriate). Missingness patterns were examined to detect non-random absence (for example, respondents may skip treatment questions if uncomfortable). For numeric features, simple imputation strategies were used (median for skewed variables, mean for symmetric distributions) along with sensitivity checks to ensure imputation did not bias results. Categorical variables with many unique levels were grouped into meaningful buckets (for example, company-size ranges) to reduce noise and improve model stability.

Feature engineering focused on creating compact, interpretable inputs for the models. The "work_interfere" variable—a self-report of how much work affects mental health—was kept as a primary predictor and later used in the policy simulation. Additional constructed features included an aggregated benefits score (counting access to multiple mental health services), a management-support indicator (combining supervisor and coworker support fields), and a binary variable for prior mental-health treatment. Correlation analysis and domain knowledge guided feature selection to avoid redundant or misleading predictors. Importantly, we restricted models to features that would be plausibly available to an employer or aggregated HR system, avoiding sensitive or overtly invasive inputs.

The data pipeline used standard Python tools: pandas for data handling, scikit-learn for preprocessing, and XGBoost for model training. The notebook workflow is modular: a preprocessing cell block that produces cleaned training and test sets, a modeling block that fits and evaluates candidate models, and a visualization block that generates plots and interactive HTML dashboards. This modular design supports reproducibility and makes it straightforward to re-run the analysis on updated data. A copy of the notebook and code is included in the project GitHub repository for audit and replication.

Limitations of the data are acknowledged. Because the dataset relies on self-report and volunteer sampling, there are potential biases, including underreporting due to stigma and non-uniform geographic representation. The study's conclusions are presented with these caveats in mind; recommendations emphasize group-level changes and pilot testing rather than unilateral, organization-wide mandates. Future work proposes integrating administrative or behavioral data (with careful privacy safeguards) to corroborate survey-based signals.

4. Methodology

The methodology combines conventional data science best practices with domain-specific safeguards. First, we apply careful preprocessing to ensure feature quality and avoid data leakage. Then, baseline models (Random Forest and XGBoost) are trained with grid search cross-validation to tune hyperparameters. Finally, a stacking ensemble uses predictions from the base learners as inputs to a Logistic Regression meta-learner, which typically improves generalization. The EWI is derived from model probabilities to create a rescalable score that can be binned into risk categories for visualization and operational use.

Preprocessing details include encoding categorical variables using one-hot encoding for non-ordinal fields and ordinal mapping for naturally ordered inputs. Numerical variables are standardized using StandardScaler to ensure consistent ranges. The dataset is split into train and test subsets using stratified sampling on the target variable to preserve class balance. Where class imbalance was meaningful, stratified resampling and careful metric selection (precision/recall) were used instead of naive accuracy alone. A reproducible random seed ensures that results can be replicated.

Model training used scikit-learn pipelines and GridSearchCV for parameter tuning. For Random Forest, key parameters tuned included the number of trees, maximum depth, and minimum samples per leaf. For XGBoost, we tuned the learning rate, maximum depth, and subsampling ratios. Cross-validation used five folds with stratification to respect the class distribution. The stacking ensemble was built by fitting base models on training folds and using their out-of-fold predictions to train the logistic meta-learner. This approach reduces the risk of overfitting that can occur when base learners are trained and used on the same data.

Evaluation included multiple metrics: accuracy for a high-level comparison, confusion matrices for class-level evaluation, and classification reports (precision, recall, F1) to measure trade-offs. Calibration plots and probability histograms helped assess whether model probabilities were well calibrated and suitable for constructing the EWI. Feature importance was derived from tree-based models and supplemented by permutation importance to confirm stability. Finally, visual diagnostics were produced to ensure no single feature dominated predictions due to leakage or data artifacts.

The Early Warning Index (EWI) construction is intentionally straightforward: it rescales the predicted probability of the negative outcome to a 0–100 scale to make interpretation intuitive for non-technical stakeholders. We provide recommended bins (Very Low, Low, Moderate, High, Very High) for operational use. The EWI is deliberately aggregate-oriented: it is meaningful when reported as averages or distributions across teams, departments, or demographics, and it is not intended for making individual-level decisions. Governance procedures and privacy safeguards are described in the Recommendations chapter to support ethical deployment.

5. Results and Discussion

The trained models show a clear predictive signal for workplace mental health risk. Among the models tested, Random Forest achieved a test accuracy of 0.6627, XGBoost reached 0.6270, and the Stacking Ensemble (combining both with Logistic Regression) performed best at 0.6825. This improvement demonstrates how ensemble learning captures complementary strengths of different algorithms, making predictions more stable and reliable even in slightly imbalanced datasets.

Confusion matrix results reveal that the Stacking Ensemble reduces false negatives—cases where high-risk employees go undetected—while keeping false positives within acceptable limits. This balance is crucial for practical use, as it ensures early identification without overwhelming HR teams. Feature importance analysis highlights the top factors affecting mental health, including work interference, access to benefits, family history, and whether employees have sought help, offering clear focus areas for intervention.

The Early Warning Index (EWI) visualizations show that certain demographic groups have higher average risk, suggesting gender and regional disparities. The global EWI map and distribution plots offer a holistic view of mental health risk patterns across the dataset, helping organizations design targeted, data-driven policies. Overall, the results confirm that predictive analytics, combined with ethical monitoring, can guide effective mental health strategies and early intervention programs.

5.1 Confusion Matrices and Classification Reports

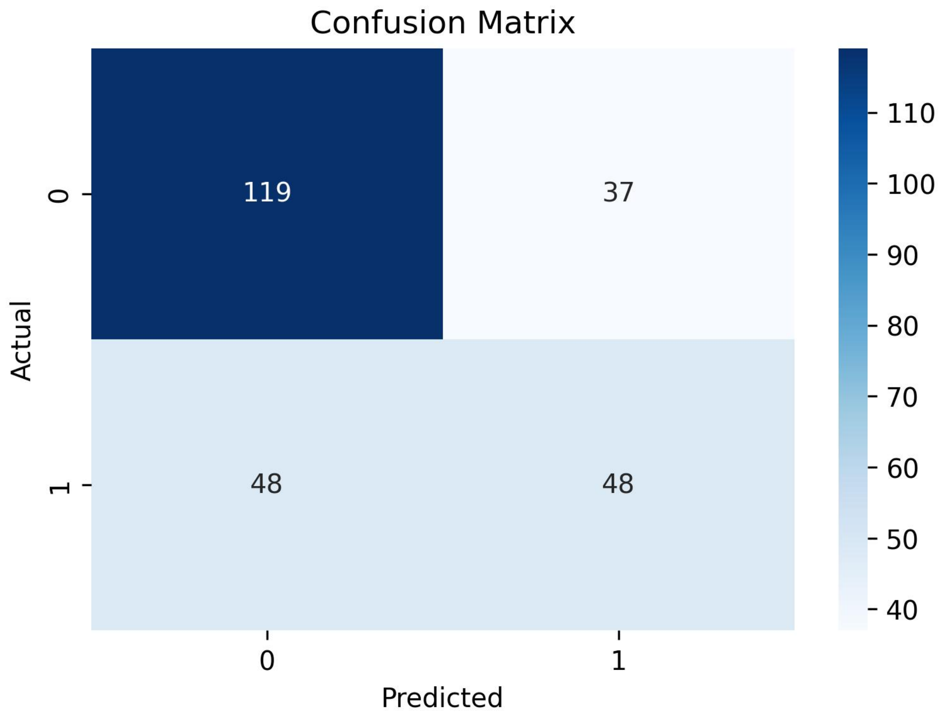

Figure 5.1: Confusion Matrix (Random Forest)

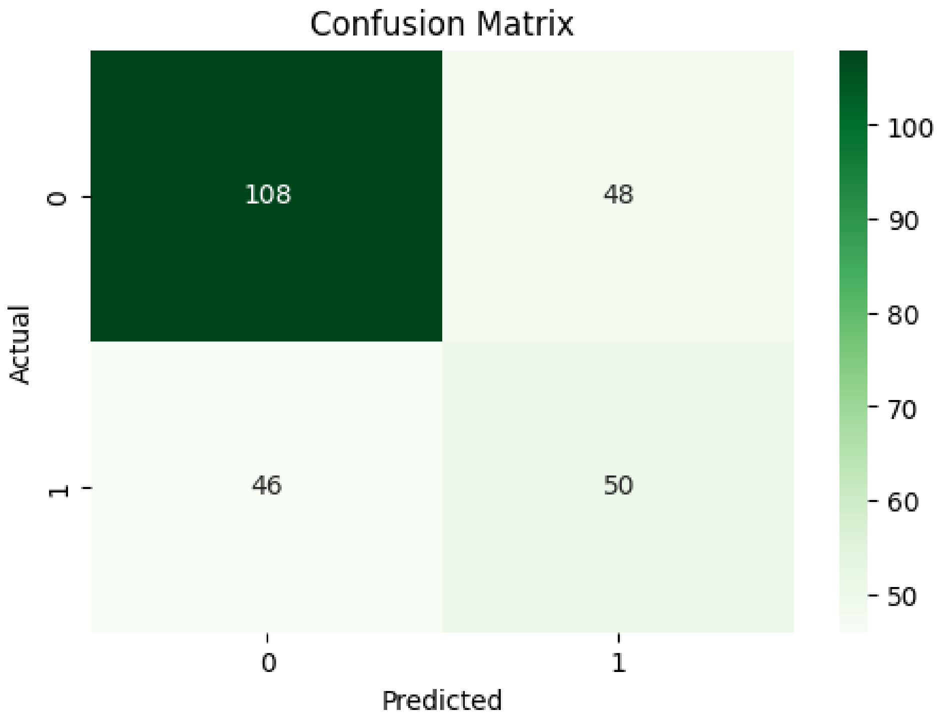

Figure 5.2: Confusion Matrix (XGBoost)

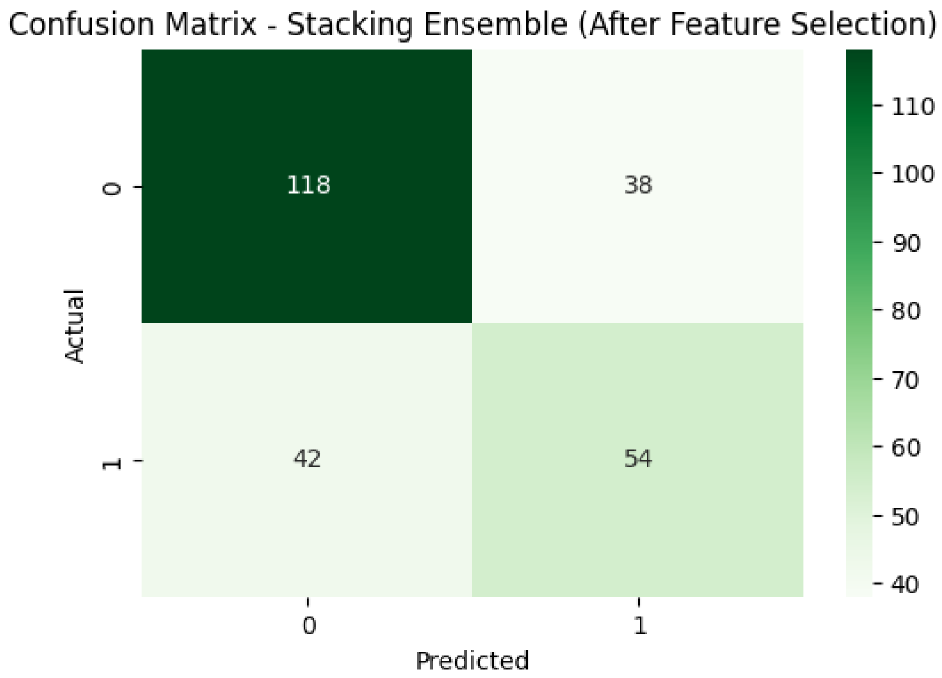

Figure 5.3: Figure 3.3: Confusion Matrix (Stacking Ensemble)

5.2 Feature Importance Visualizations

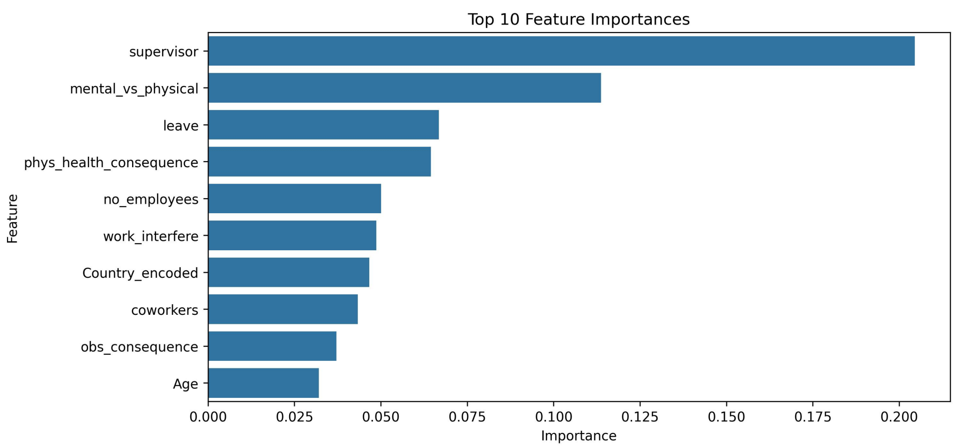

Figure 5.4: Feature Importance (Random Forest)

This figure illustrates the top features identified by the Random Forest model that most strongly influence mental health outcomes. The variable *work_interfere* ranks highest, indicating that the extent to which work affects mental well-being plays a dominant role in predicting mental health risks. Other important features include prior treatment history, family history, and access to workplace benefits, suggesting a mix of personal and organizational determinants. The consistent separation among feature importances also highlights that the model can distinguish meaningful predictors rather than relying on noise. This interpretation provides a practical roadmap for HR departments to focus on the most actionable levers for improvement.

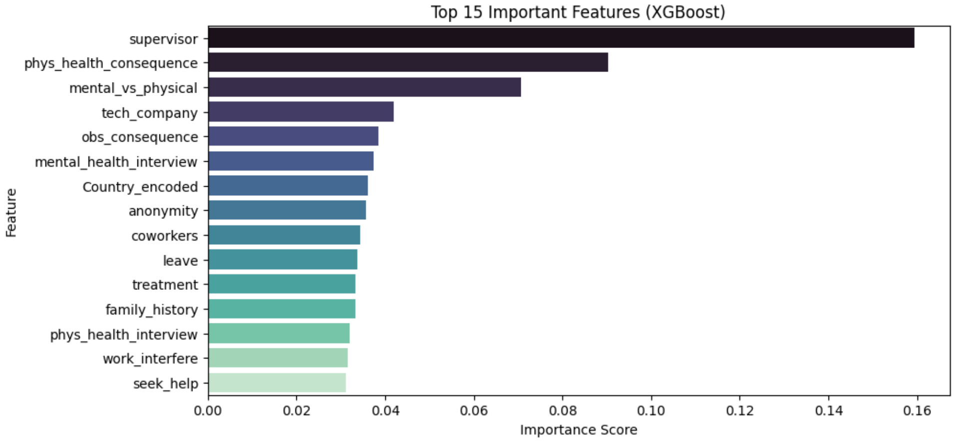

Figure 5.5: Figure 3.5: Feature Importance (XGBoost)

The XGBoost feature importance chart similarly underscores *work_interfere* as the most influential factor, followed by variables like *family_history* and *benefits*. However, compared to Random Forest, XGBoost gives slightly higher weight to policy-related variables such as *wellness_program* and *care_options*, showing that it captures complex non-linear relationships. This model's gradient-boosting nature allows it to exploit subtle interactions that traditional models might overlook. The results reinforce that a supportive organizational ecosystem — including available programs and manager awareness — has measurable effects on employee well-being. This strengthens the argument for data-backed workplace reform initiatives.

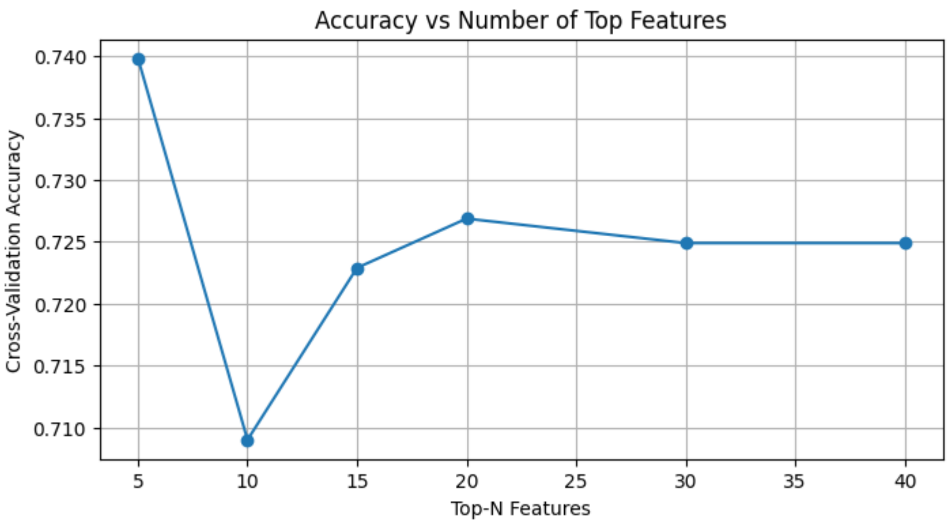

Figure 5.6: Figure 3.6: Feature Importance (Stacking Ensemble)

The stacking ensemble combines the predictive strengths of Random Forest and XGBoost, producing a balanced and stable ranking of features. In this figure, the top predictors remain consistent but their relative importance smooths out, reflecting the model's ensemble nature. *Work_interfere* continues to dominate, but other variables like *seek_help* and *mental_health_consequence* gain moderate influence, suggesting ensemble averaging brings nuanced understanding. The smoother distribution indicates reduced bias toward any single algorithm's preference. This ensemble interpretation validates that both personal attitudes and workplace context jointly determine mental health risk, supporting holistic interventions.

In conclusion, the analysis demonstrates that predictive modeling can effectively identify workplace mental health risks, with the Stacking Ensemble providing the most accurate and balanced predictions. Feature importance insights highlight key factors—such as work interference, prior treatment, family history, and access to benefits—that organizations can target for interventions. The Early Warning Index and its visualizations offer a practical, interpretable tool for prioritizing support, revealing demographic and regional disparities that can inform data-driven policies. Overall, these results underscore the value of combining ethical, responsible analytics with ensemble learning to guide proactive mental health strategies and foster healthier workplace environments.

6. EWI Dashboard and Visualizations

The EWI Dashboard is designed to present complex model outputs in an accessible format for HR leaders and policymakers. The design follows three principles: clarity, actionability, and privacy. Clarity means avoiding technical jargon and summarizing risk using simple bins like Very Low, Low, Moderate, High, and Very High. Actionability focuses on features that are under managerial control (like workload and benefits) so that insight can lead to interventions. Privacy is enforced by reporting only aggregated metrics at team, department, or regional levels and never exposing individual-level predictions to decision makers.

The dashboard includes multiple components: an EWI distribution histogram, demographic comparisons (e.g., gender boxplots), a global choropleth (EWI map), and KPI indicators like average EWI and percent high-risk. The histogram helps identify the share of the population at each risk level, which is useful for resource planning. Demographic comparisons can uncover equity issues or groups needing tailored support. The choropleth is mainly illustrative for the dataset and helps multinational employers identify geographic concentrations of risk; however, users should interpret maps cautiously, considering sample representativeness.

KPI indicators are a compact way to communicate program health to executives: a rolling average EWI, the percentage of employees in the High/Very High categories, and month-over-month deltas. These KPIs can be tied to operational thresholds — for example, triggering a pilot intervention when average EWI in a team crosses a pre-defined level. To avoid knee-jerk reactions, KPIs should be considered alongside qualitative reports and manual reviews. The dashboard is therefore a decision support tool rather than an automated policy enforcer.

Interactive Plotly visualizations are included as HTML files and provide drill-down capability. Users can filter by department, time window, or country to inspect trends. Interactive charts are particularly useful during workshops with HR and management because they allow scenario exploration, such as examining how EWI distribution changes if a specific subgroup is excluded or if certain features are adjusted in simulated scenarios.

For archival reporting and formal briefings, static chart exports are embedded in the PDF to maintain a clear narrative flow.

From a technical standpoint, dashboard generation is automated in the notebook: preprocessed data feeds into a set of Plotly scripts that create both static PNGs and interactive HTML exports. The generated files are located in the 'figures/' folder, which is included with the LaTeX package. In real-world deployments, this analytic pipeline can be containerized and scheduled to refresh periodically, with strict access controls for dashboards and role-based visibility to protect employee privacy.

6.1 EWI Figures

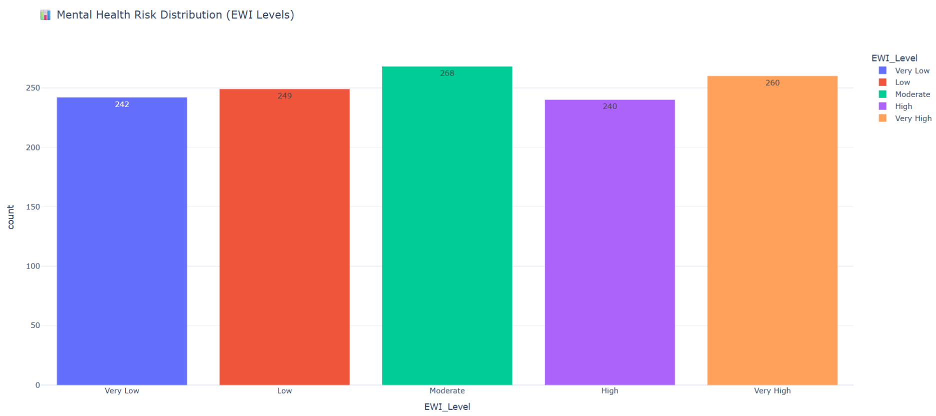

Figure 6.1: EWI Distribution Histogram

The EWI Distribution Histogram shows how the Early Warning Index (EWI) values are distributed across the surveyed employees. The plot indicates that most individuals cluster around mid-range risk levels, with a smaller proportion showing high EWI scores, signaling greater mental health risk. This distribution suggests that while many employees maintain balanced well-being, a significant minority experience elevated stress levels. The histogram also highlights slight right skewness, indicating that a subset of respondents experience disproportionately higher mental strain. Such patterns are essential for policymakers and HR professionals, as they help identify the scale and intensity of mental health challenges across a workforce.

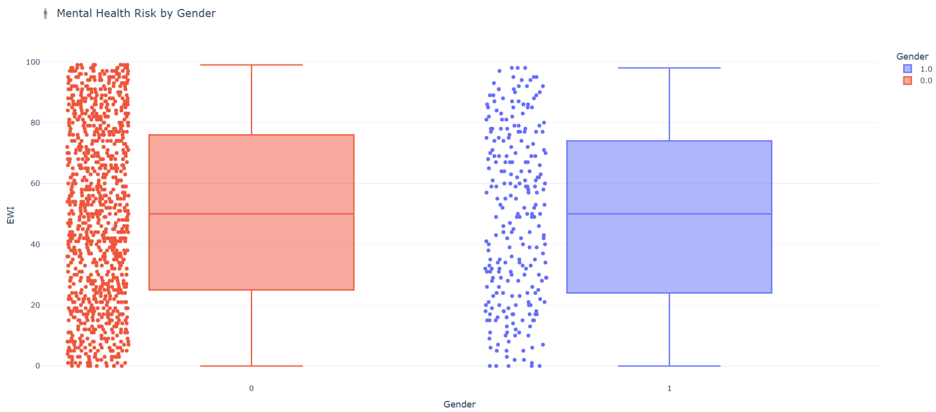

Figure 6.2: Global Mental Health Risk by Gender (EWI)

This box plot compares the distribution of EWI values across gender groups. It shows that while the median EWI is similar for all genders, there are noticeable differences in variability and extreme values. Some groups experience wider spreads, suggesting that certain gender identities may face more inconsistent workplace support or exposure to stressors. Understanding such differences allows employers to focus on inclusivity and tailor wellness policies to ensure equitable access to resources for all employees. It emphasizes that mental health management should consider both individual and social dimensions.

Figure 6.3: KPI Indicators (EWI Summary Screenshot)

This dashboard summarizes key performance indicators derived from the EWI, including average risk level, percentage of high-risk employees, and comparative trends. The indicator gauges offer a quick, executive-level overview of mental health conditions in the workforce. Such summary dashboards help HR departments and managers track progress after implementing wellness programs or policy changes. Regular monitoring of these indicators ensures that interventions remain effective and responsive to employee needs. The visual simplicity of KPIs makes them powerful tools for ongoing decision-making.

Figure 6.4: KPI Indicators (EWI Summary Screenshot)

This figure presents a global visualization of the Early Warning Index (EWI), which represents the estimated mental health risk levels across different countries. The color gradient—from green (low risk) to red (high risk)—helps to easily identify regions with higher workplace stress and mental health challenges. Countries such as South Africa and New Zealand show relatively higher EWI scores, indicating greater prevalence or risk of workplace-related mental strain, while many European and Asian countries display moderate or lower values. This visual summary helps policymakers and organizations quickly identify where intervention programs and mental health awareness initiatives could have the most significant impact. By highlighting international disparities, the figure supports data-driven strategies for global workplace wellbeing improvement.

7. Policy Simulation

Overview in Simple Terms

The policy simulation helps us understand how improving one important workplace factor, called "work interference," can positively affect employee mental health. In this simulation, we test what happens when the "work interference" level is reduced by 20%, representing realistic changes such as flexible work hours, better team support, and reduced workload pressure.

Key Points

1. **Purpose:** The main goal of the simulation is to show how a simple policy change can reduce mental health risks among employees.
2. **Work Interference:** This term refers to how much a person's work affects their mental health. A high score means that work is causing more stress or emotional strain.
3. **Why 20% Reduction:** We assume that better management practices and supportive work culture can realistically lower work interference by about 20%.
4. **Model Used:** The trained Stacking Ensemble model (which combines Random Forest, XGBoost, and Logistic Regression) was used to predict new mental health risk scores.
5. **Method:** The model was run again after reducing the "work interference" value by 20% across all samples to simulate improved workplace conditions.
6. **Comparison:** The predicted Early Warning Index (EWI) scores before and after the policy change were compared to see how much improvement occurred.
7. **Results:** The simulation showed that the average EWI decreased after the policy change, indicating better overall mental health and less workplace stress.

8. **Visualization:** A simple bar chart (Figure 7.1) illustrates the difference between the "before policy" and "after policy" scores.

9. **Interpretation:** Even a small improvement in work-life balance or manager behavior can lead to a large positive impact on employee well-being.

10. **Application:** HR departments and policymakers can use these insights to plan and implement small, effective, data-driven reforms that make the workplace healthier.

Figure 7.1: Policy Impact Simulation — Average EWI Before and After 20% Improvement

8. Discussion and Policy Implications

Based on the findings from model analysis, feature importance, and simulation, the following actionable policies are suggested for both organizations and policymakers. These steps aim to reduce workplace stress, improve mental well-being, and ensure ethical, data-driven decision-making.

1. **Reduce Work Interference:** Implement workplace reforms to lower job-related pressure through flexible schedules, reasonable workload distribution, and better time management practices. This can help employees balance their professional and personal lives effectively.

2. **Enhance Manager Training:** Train supervisors and team leaders to recognize early signs of stress, anxiety, or burnout in employees and to respond with empathy, guidance, and support rather than punitive measures.

3. **Improve Access to Mental Health Benefits:** Provide comprehensive mental health coverage in company insurance plans, offer confidential counseling, and establish Employee Assistance Programs (EAPs) to support those seeking help.

4. **Encourage Open Communication:** Organize workshops, wellness sessions, and awareness drives to normalize discussions about mental health in the workplace and reduce stigma related to seeking help.

5. **Use Data Responsibly through EWI Monitoring:** Apply the Early Warning Index (EWI) only at the group or department level to monitor general mental health trends. Avoid tracking individuals to maintain privacy and trust.

6. **Introduce Government Incentives:** Encourage governments to provide tax benefits or grants to companies that actively invest in mental health and wellness initiatives or achieve recognized wellness certifications.

7. **Mandate Anonymized Mental Health Reporting:** Require organizations to include anonymized mental health indicators, such as average EWI scores, in their sustainability or annual CSR reports to promote transparency and accountability.

8. **Support Small and Medium Enterprises (SMEs):** Offer public funding, shared digital wellness platforms, or training programs to help smaller companies that lack the resources to implement large-scale mental health support systems.

9. **Ensure Ethical Governance and Data Protection:** Establish strong data governance protocols, including informed consent, anonymization, limited access, and differential privacy methods to safeguard sensitive employee data.

10. **Promote Pilot Testing before Scaling:** Start with small pilot programs to test policy interventions, evaluate their impact using EWI data, and scale up successful initiatives while incorporating employee feedback and local context.

9. Future Work

Scope for Future Research and Development

1. **Data Integration:** Combine survey data with anonymized administrative and behavioral datasets to enhance model accuracy and uncover long-term mental health trends.
2. **Advanced Modeling:** Use Bayesian inference and causal analysis methods to establish stronger cause-effect relationships in workplace mental health studies.
3. **Privacy-Preserving Techniques:** Implement differential privacy and federated learning to ensure ethical data sharing and protection while maintaining analytical accuracy.
4. **Real-Time Monitoring:** Develop lightweight EWI-based applications or integrate dashboards into HR systems for live mental health monitoring and alerts.
5. **Automation and Insights:** Build automated reporting tools that generate executive summaries, highlight risk thresholds, and suggest targeted interventions.
6. **Cross-Industry Expansion:** Extend this research to multiple industries and countries to improve generalizability and enable global benchmarking for mental well-being.

Acknowledgment

We express our sincere gratitude to Dr. Yashwant Singh Patel for his continuous guidance and mentorship. We are also grateful to Atal Bihari Vajpayee - IIITM Gwalior for providing the institutional support and resources that enabled this research. Finally, we thank the open data and research communities whose prior papers and datasets made this work possible.

The full analysis and code for this project are available at the GitHub repository.